



Host



Kernel



```

void kernel_main(kernel_arg_t* arg) {
    auto smem_base = (uint8_t*)__local_mem();
    auto B_smem = (ctx::input_t*)
        (smem_base + smem_b_off);
    ctx::fragment_acc fragC;
    ctx::fill_fragment(fragC, 0);

    for (uint32_t k = 0; k < K; k += ctx::tileK) {
        auto A_warp = (ctx::input_t*)
            (smem_base + warp_rank * per_warp_section);
        // Cooperatively stage A' and B in smem
        for (...) A_warp[i] = pA[...];
        for (...) B_smem[i] = pB[...];

        auto meta_sp = pMetaSp +
            (tile_row_idx_w * num_k_tiles + k_tile)
            * meta_words_per_tile;
        auto desc_b = vt::vx_make_smem_desc(
            B_smem, ctx::tileK * sizeof(ctx::input_t));

        ctx::fragment_a fragA;
        ctx::load_sp_metadata(fragA, meta_sp);
        #ifdef WGMMMA_RS
        ctx::load_matrix_sync(
            fragA, A_warp, ctx::tileK / 2);
        ctx::wgmma_sync(fragC, fragA, desc_b, fragC);
        #else
        auto desc_a = vt::vx_make_smem_desc(
            A_warp, (ctx::tileK / 2) * sizeof(ctx::input_t));
        ctx::wgmma_sync(fragC, desc_a, desc_b, fragC);
        #endif
    }
    ctx::store_matrix_sync(pTileC, fragC, N);
}
  
```